

# **Intro to HW Design & Externs for P4→NetFPGA**

**CS344 – Lecture 5**

# Announcements

---

- **Wednesday:**

- Static router interoperability pretesting
- Group meetings with instructors (Gates 315)

- **Thursday:**

- Design challenge project proposals

# Outline


---

- **Intro to HW design & FPGAs**
- **How to build your own stateful extern for P4→NetFPGA**
- **Design Challenge Projects**

# Logic Gates


**AND Gate**

A	B	Y
0	0	0
0	1	0
1	0	0
1	1	1




**OR Gate**

A	B	Y
0	0	0
0	1	1
1	0	1
1	1	1



**XOR Gate**

A	B	Y
0	0	0
0	1	1
1	0	1
1	1	0



# Software vs Hardware Design

---

- **Software Design**

- Functionality as sequence of instructions for CPU
- Language: C, C++, Python, etc.

- **Hardware Design**

- Functionality as digital circuit
- Language: Verilog, VHDL

- **A simple example:**

```
if (a > b) {  
    res = a;  
}  
else {  
    res = b;  
}  
return res;
```

# A Simple Example

## Software (C)


```
if (a > b) {  
    res = a;  
}  
else {  
    res = b;  
}  
return res;
```



```
cmp a b tmp  
beqz tmp else  
store r0 a  
return  
else: store r0 b  
return
```

## Hardware (Verilog)

```
wire a, b;  
reg res;  
always @(*)  
    if (a > b) begin  
        res = a;  
    end  
    else begin  
        res = b;  
    end  
end
```



# Verilog

---

- **Basic data types:**

- **wire**

- Example: `wire [15:0] B;`
    - Used for combinational (stateless) logic only

- **reg**

- Example: `reg [7:0] A;`
    - *Can* be used to hold state

- **Example Usage:**

- `assign B = {A[7:0], A[7:0]}; // Assignment of bits`
  - `reg [31:0] Mem [0:1023]; // 1K word memory`

# Combinational vs Sequential Logic

---

- **Combinational Logic**


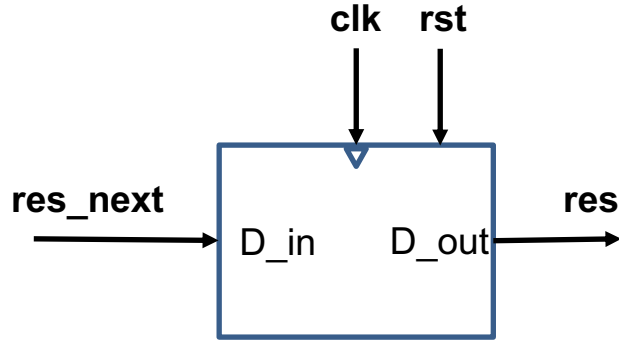
- Made of logic gates
- No memory elements
- Outputs settle to stable values after “short” logic delay

- **Sequential Logic**

- Combinational circuits and memory elements
- Used to store state
- Output depends on inputs and current state




# Adding State: Flip-Flop



```
reg res;  
always @(posedge clk)  
    if (rst)  
        res <= 0;  
    else  
        res <= res_next;
```


# Register the Output

```
wire a, b;  
reg res_next;  
always @(*)  
    if (a > b) begin  
        res_next = a;  
    end  
    else begin  
        res_next = b;  
    end  
end
```



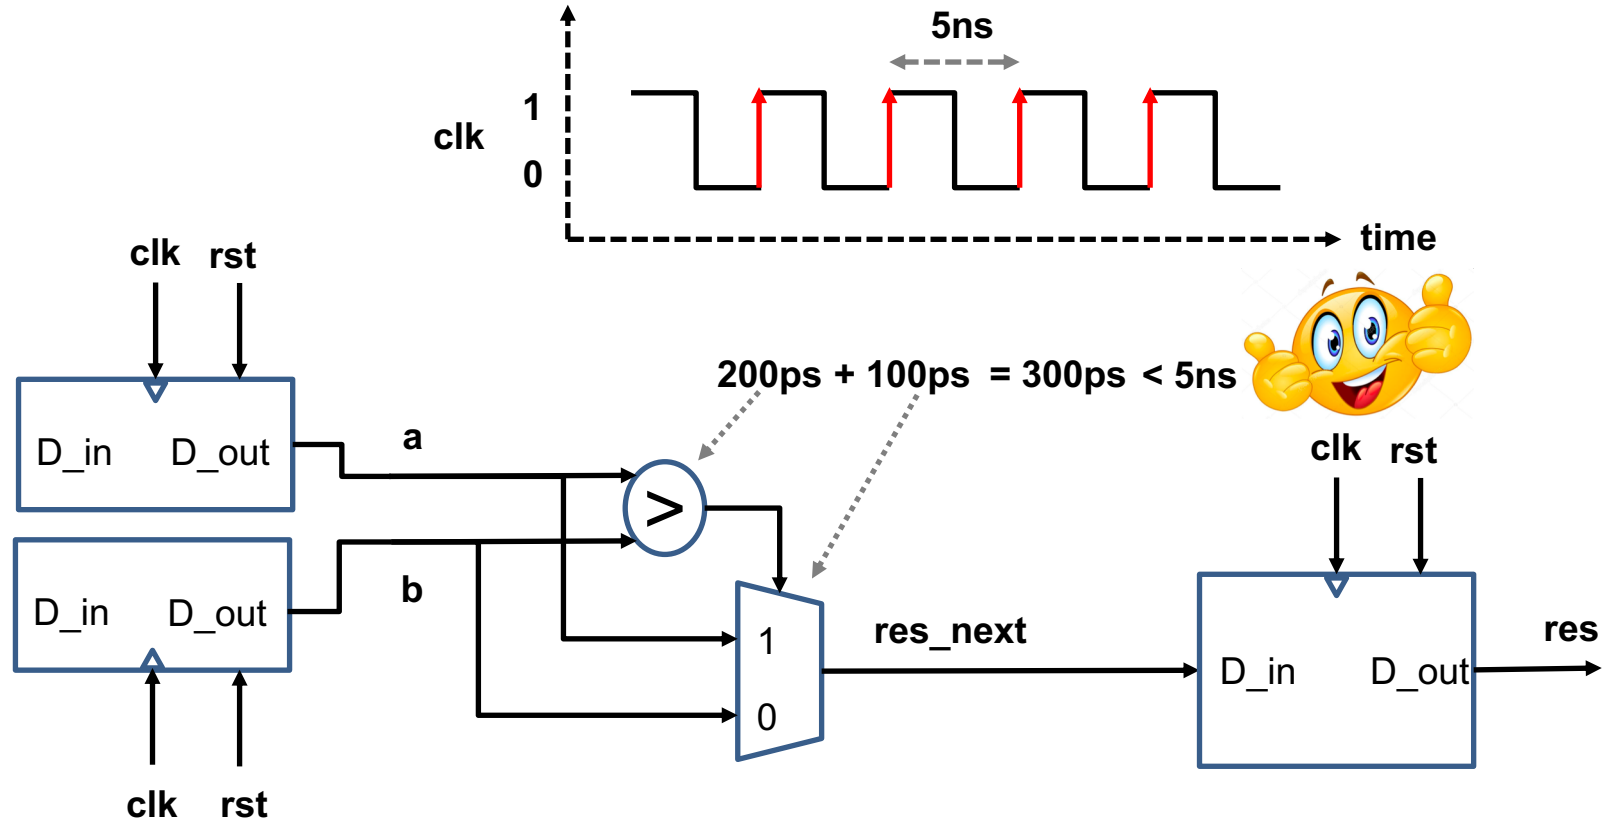
Combinational Logic

```
reg res;  
always @(posedge clk)  
    if (rst)  
        res <= 0;  
    else  
        res <= res_next;
```



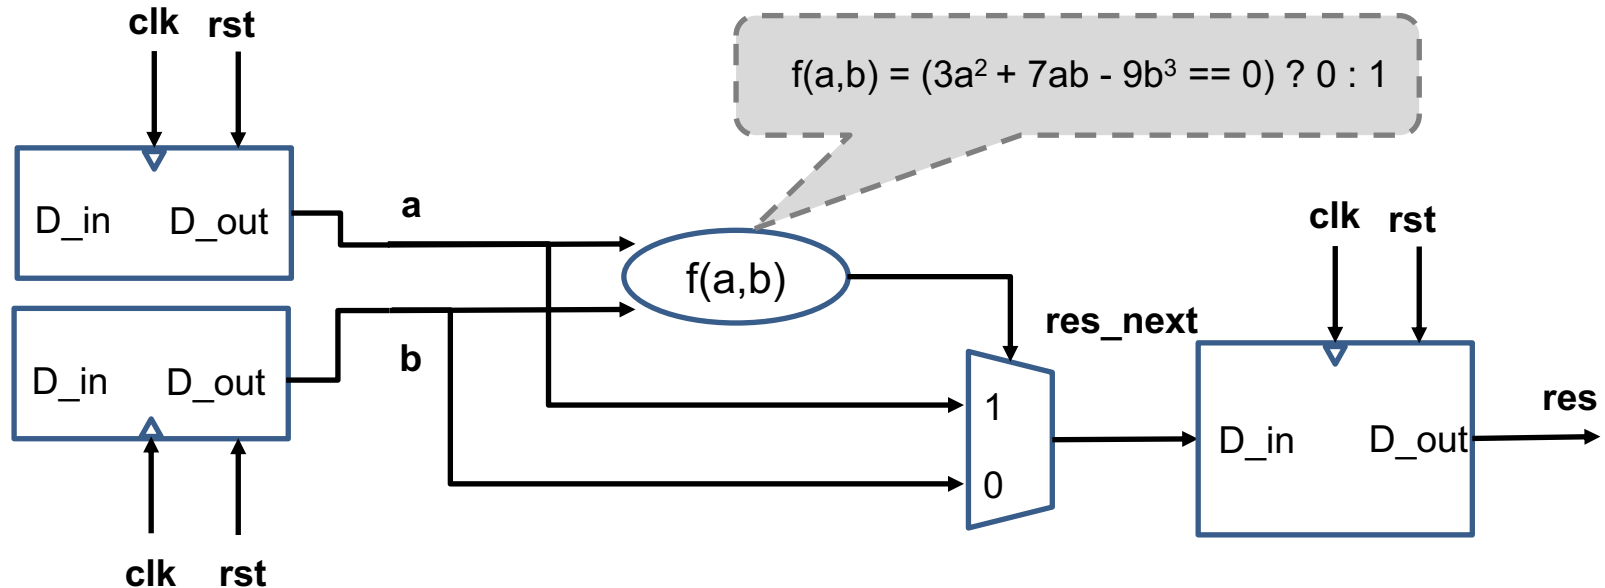

Sequential Logic

# Set Up Time Constraints

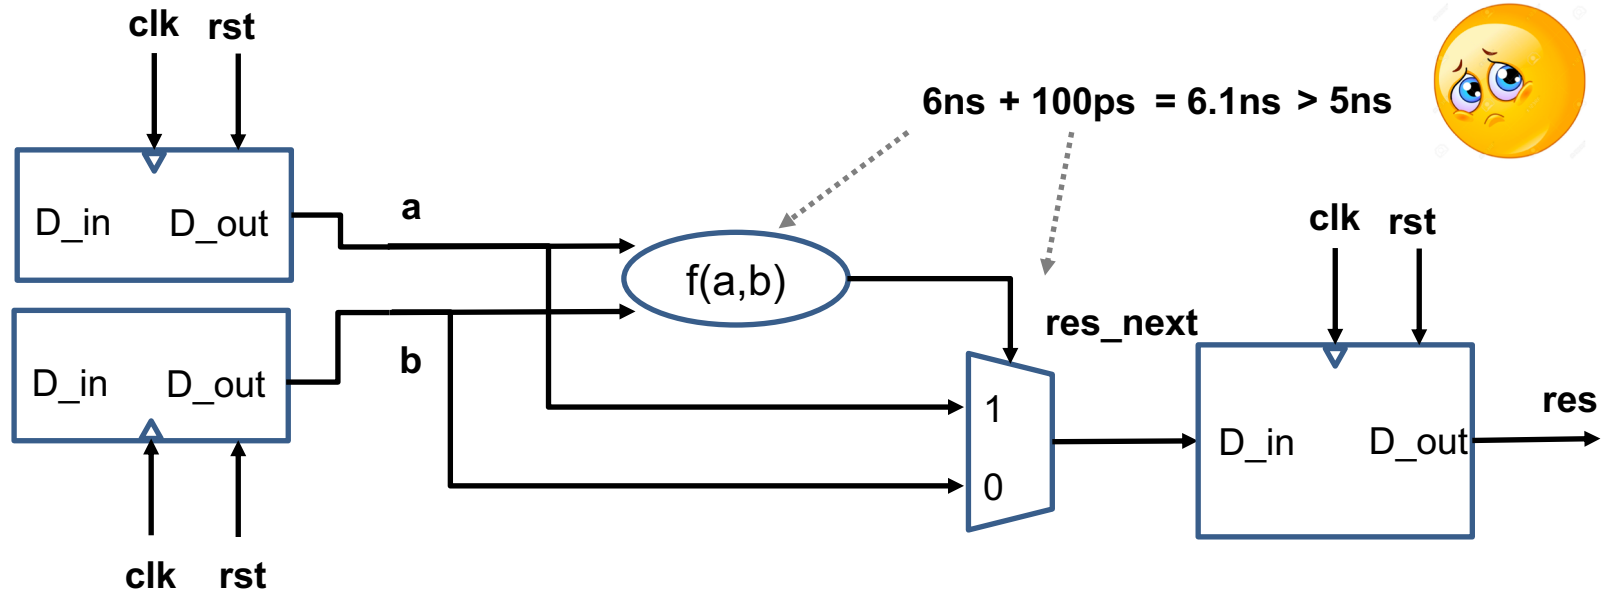



**Assumes clock is perfectly synchronized at all flip-flops!**

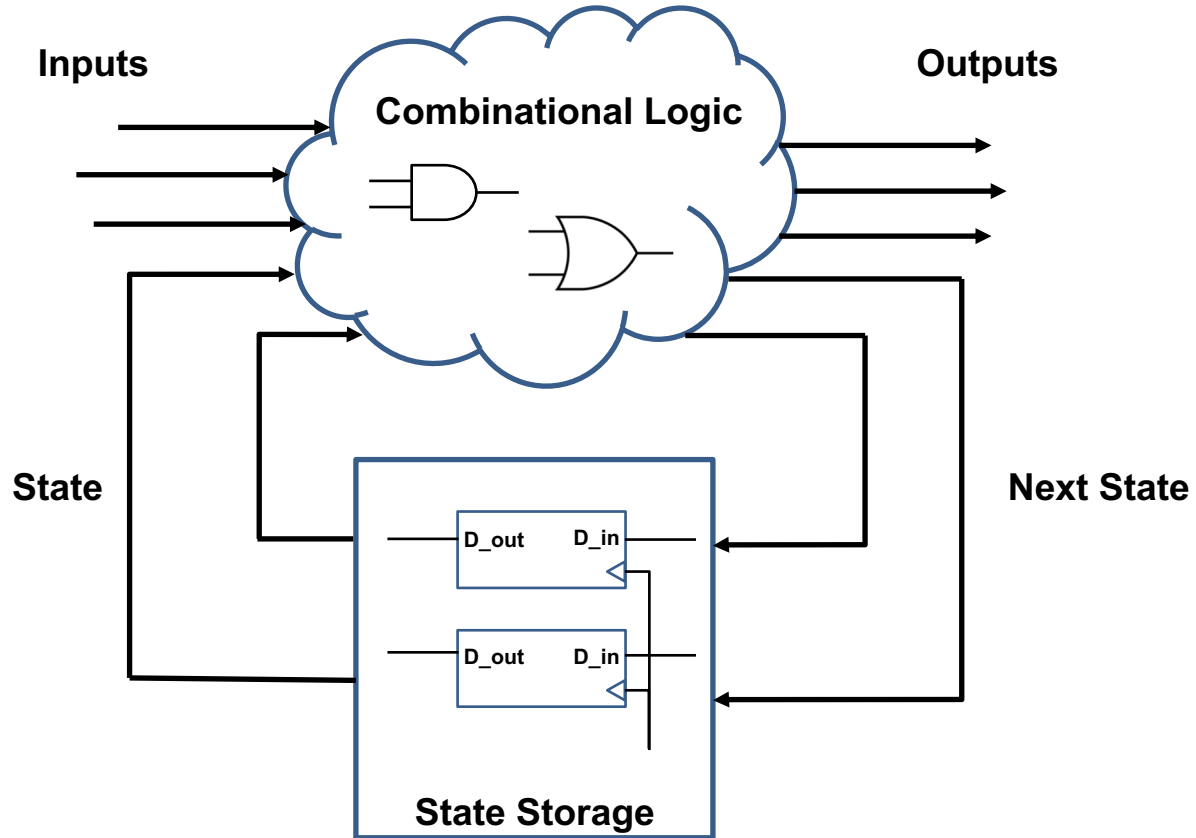
# Set Up Time Constraints



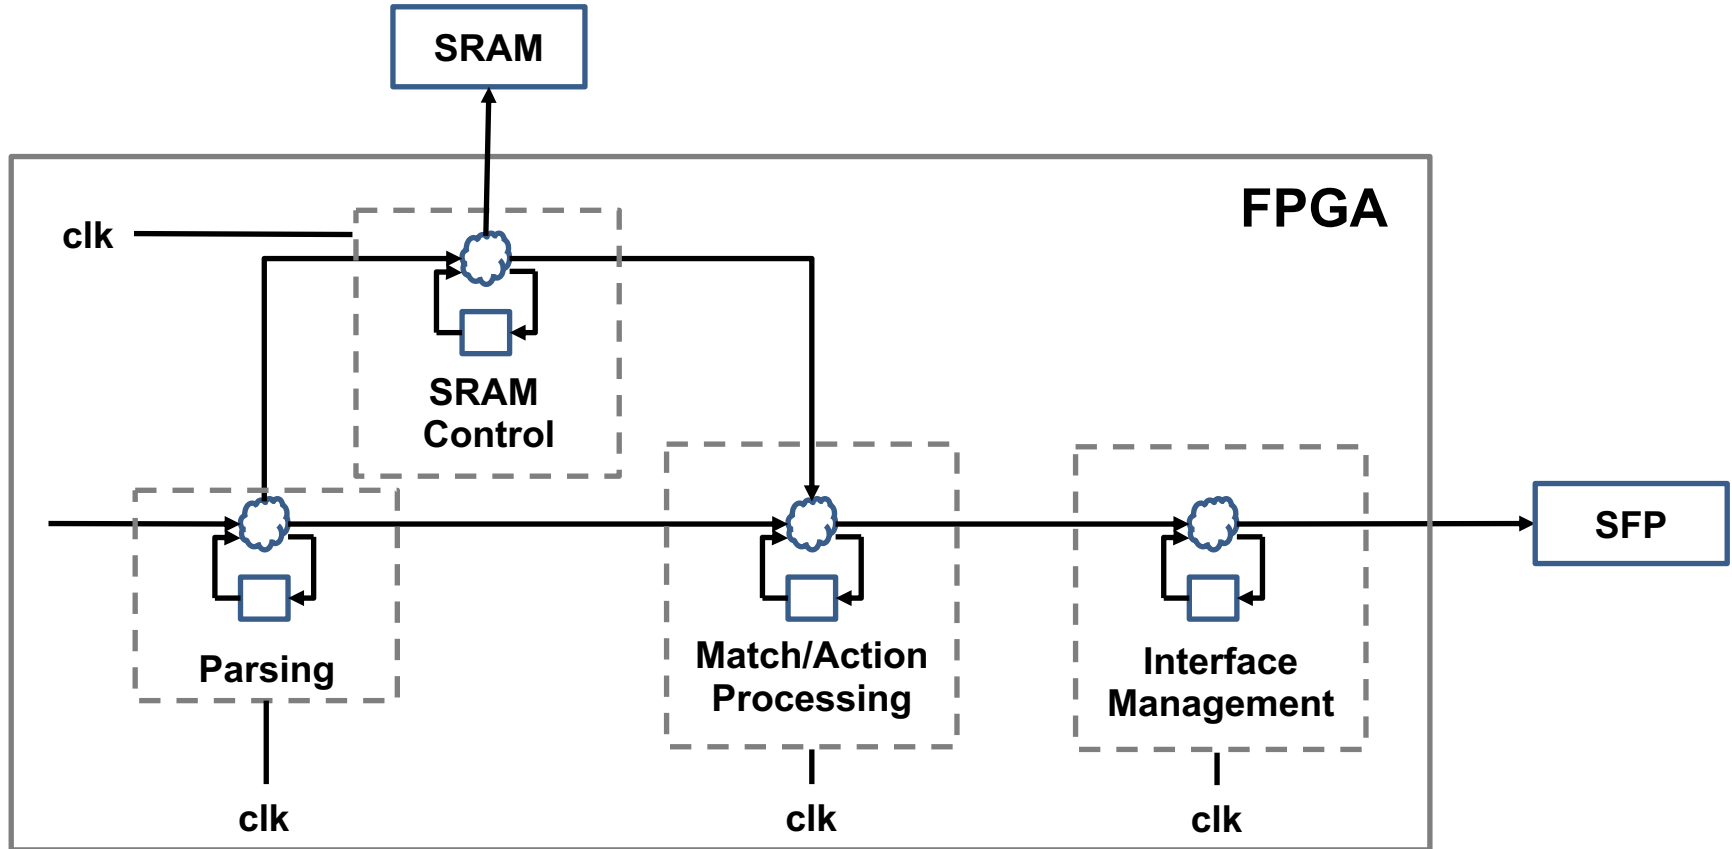
# Set Up Time Constraints



# Finite State Machine (FSM)

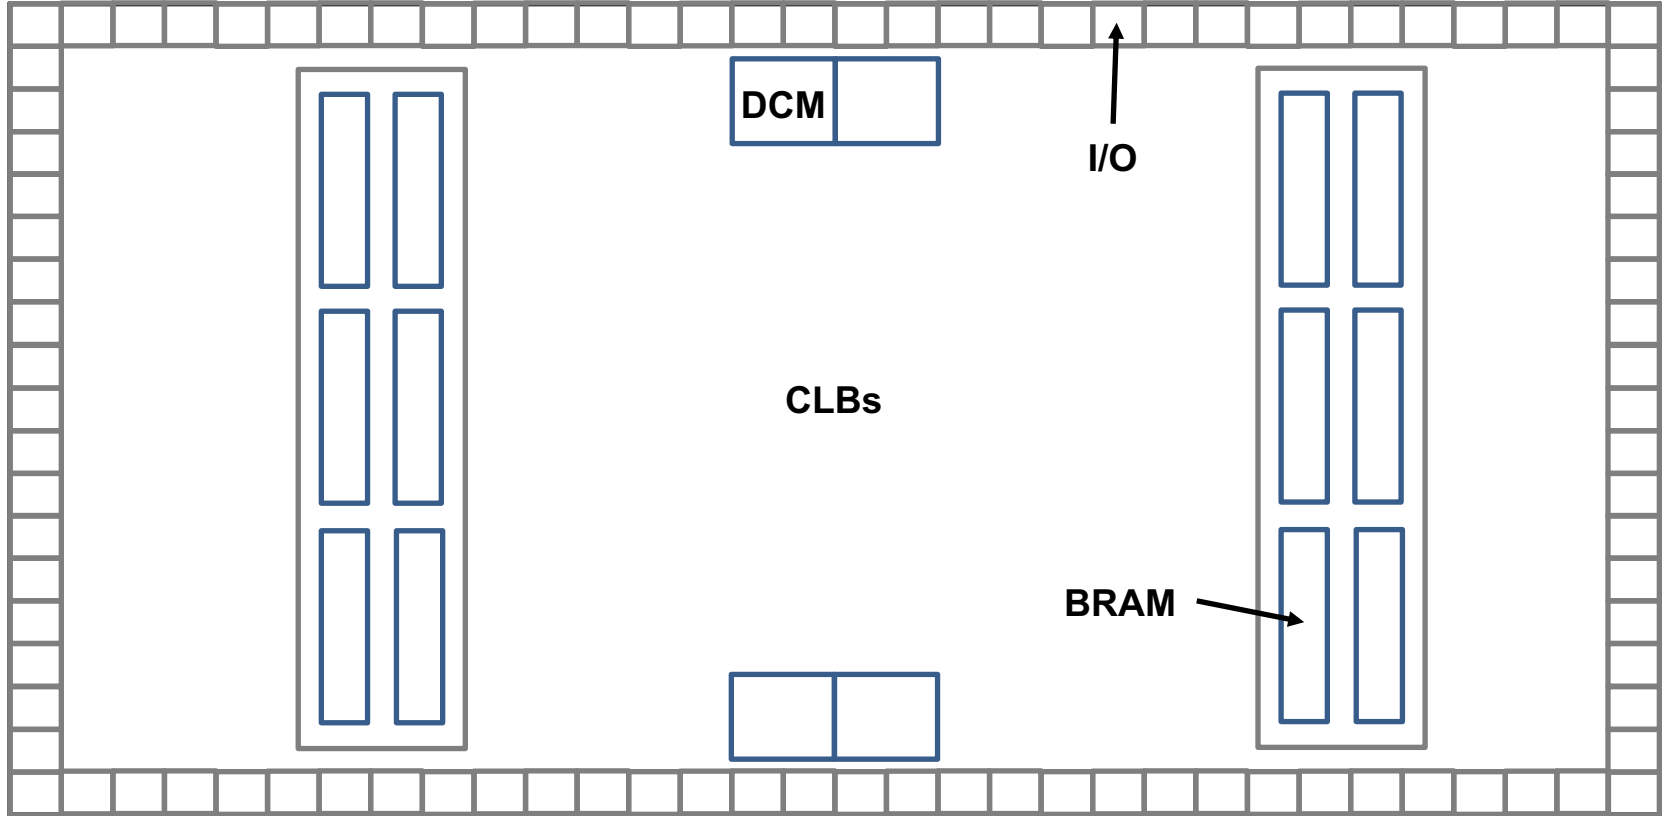


# FSMs Are Everywhere...



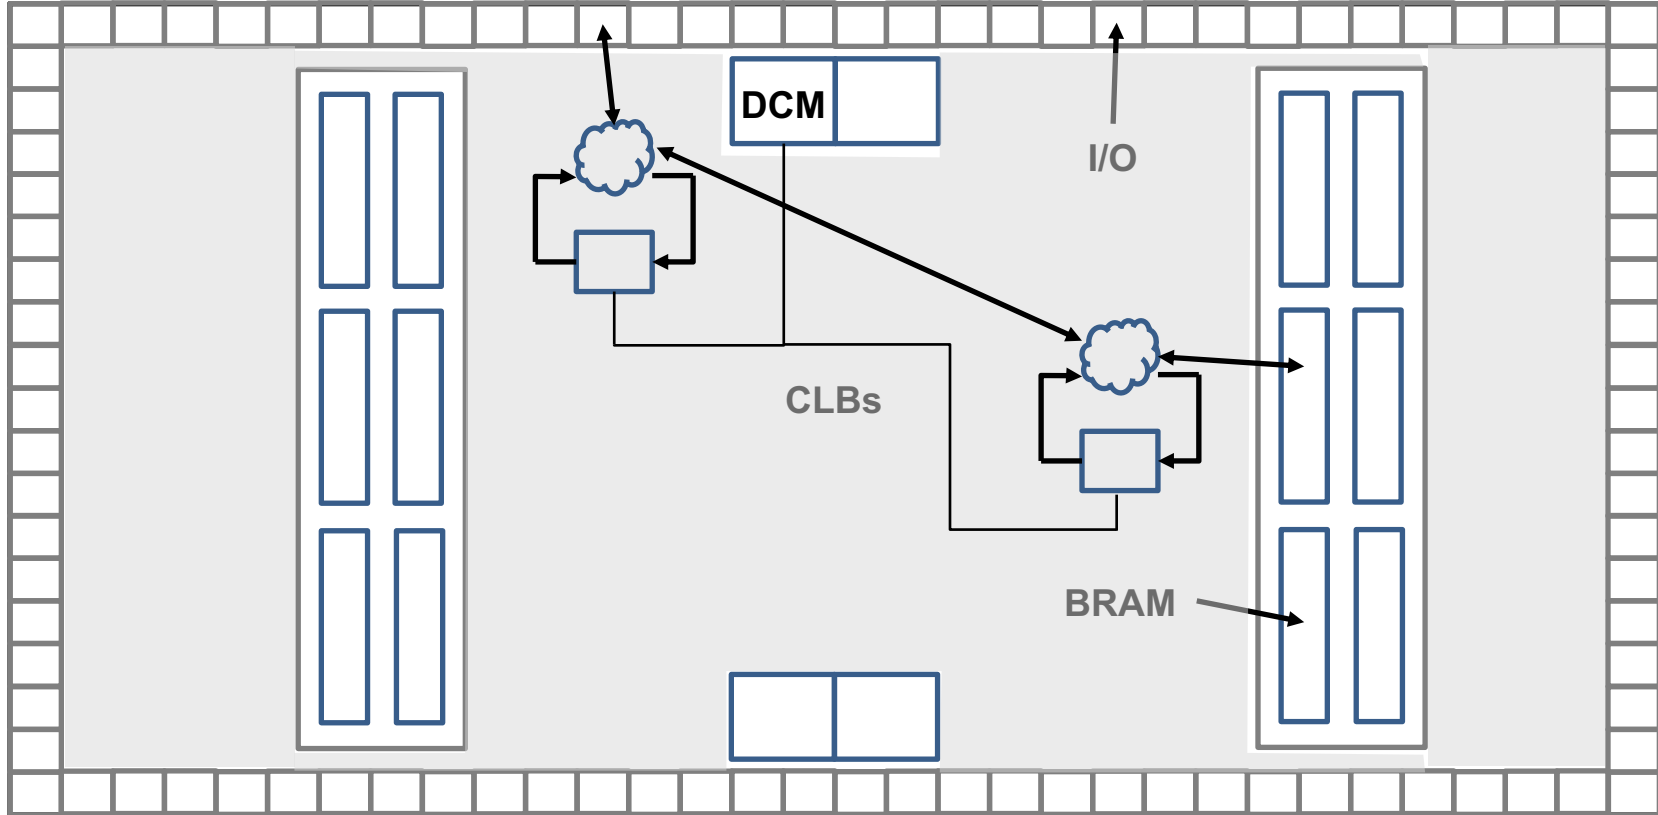
# What is an FPGA?

---





# What is an FPGA?



# FPGA Design Flow

---

- **RTL (Verilog) Design**
- **RTL Behavioral Simulation**
- ***Synthesis***
  - Decompose design into well defined logic blocks that are available on FPGA
- ***Place and Route***
  - Figure out exactly which logic blocks to use and how to route between them
- **Bitstream Generation**
- **HW Testing**

# P4→NetFPGA Extern Function library

- HDL modules invoked from within P4 programs

- Stateful Atoms [1]

Atom	Description
R/W	Read or write state
RAW	Read, add to, or overwrite state
PRAW	Predicated version of RAW
ifElseRAW	Two RAWs, one each for when predicate is true or false
Sub	IfElseRAW with stateful subtraction capability

- Stateless Externs

Atom	Description
IP Checksum	Given an IP header, compute IP checksum
LRC	Longitudinal redundancy check, simple hash function
timestamp	Generate timestamp (granularity of 5 ns)

- Add your own!

# Build a new extern: reg\_srw

---

- **Specifications:**

- Single state variable
- Can either read or write state
- Produces result in 1 clock cycle
- Not accessible by control-plane


- **P4 API:**

```
#define REG_READ 0
#define REG_WRITE 1
@Xilinx_MaxLatency(1)
@Xilinx_ControlWidth(0)
extern void <reg_name>_reg_srw<D>(in D newVal,
                                   in bit opCode,
                                   out D result);
```

# reg\_srw Next State Logic

opcodes:

- READ = 0
- WRITE = 1

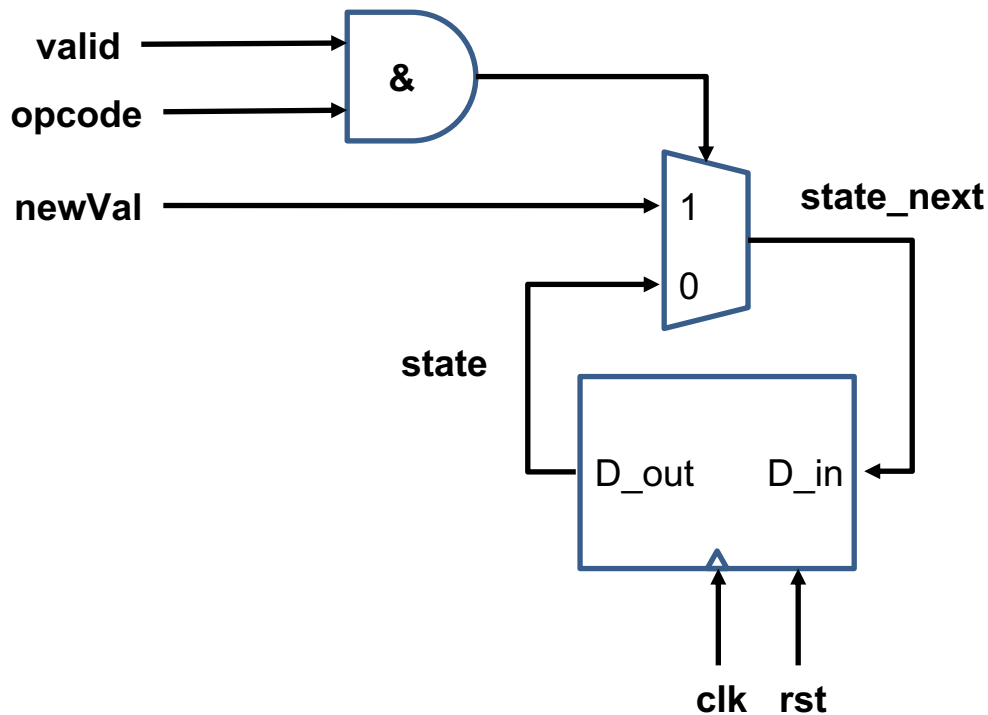


```
wire valid, opcode;
wire [7:0] newVal;
reg [7:0] state_next;
always @(*)
    if (valid & opcode)
        state_next = newVal;
    else
        state_next = state;
```

# reg\_srw Finite State Machine

opcodes:

- READ = 0
- WRITE = 1



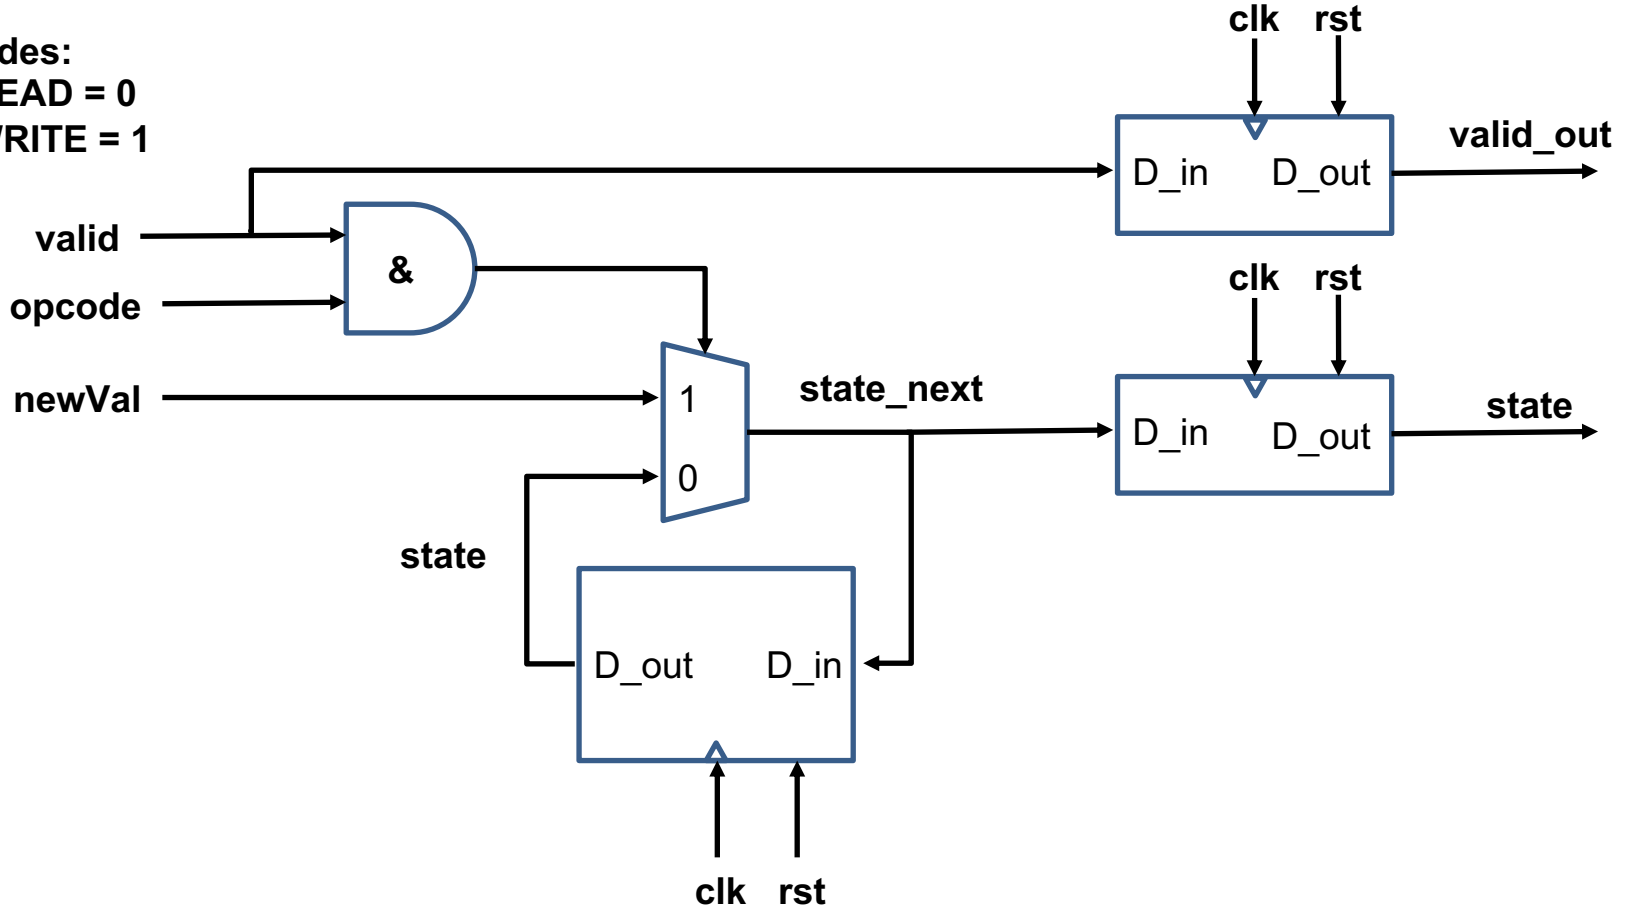
```
wire valid, opcode;
wire [7:0] newVal;
reg [7:0] state_next;
always @(*)
    if (valid & opcode)
        state_next = newVal;
    else
        state_next = state;
```

```
reg [7:0] state;
always @(posedge clk)
    if (rst)
        state <= 0;
    else
        state <= state_next;
```

# Register the Outputs

opcodes:

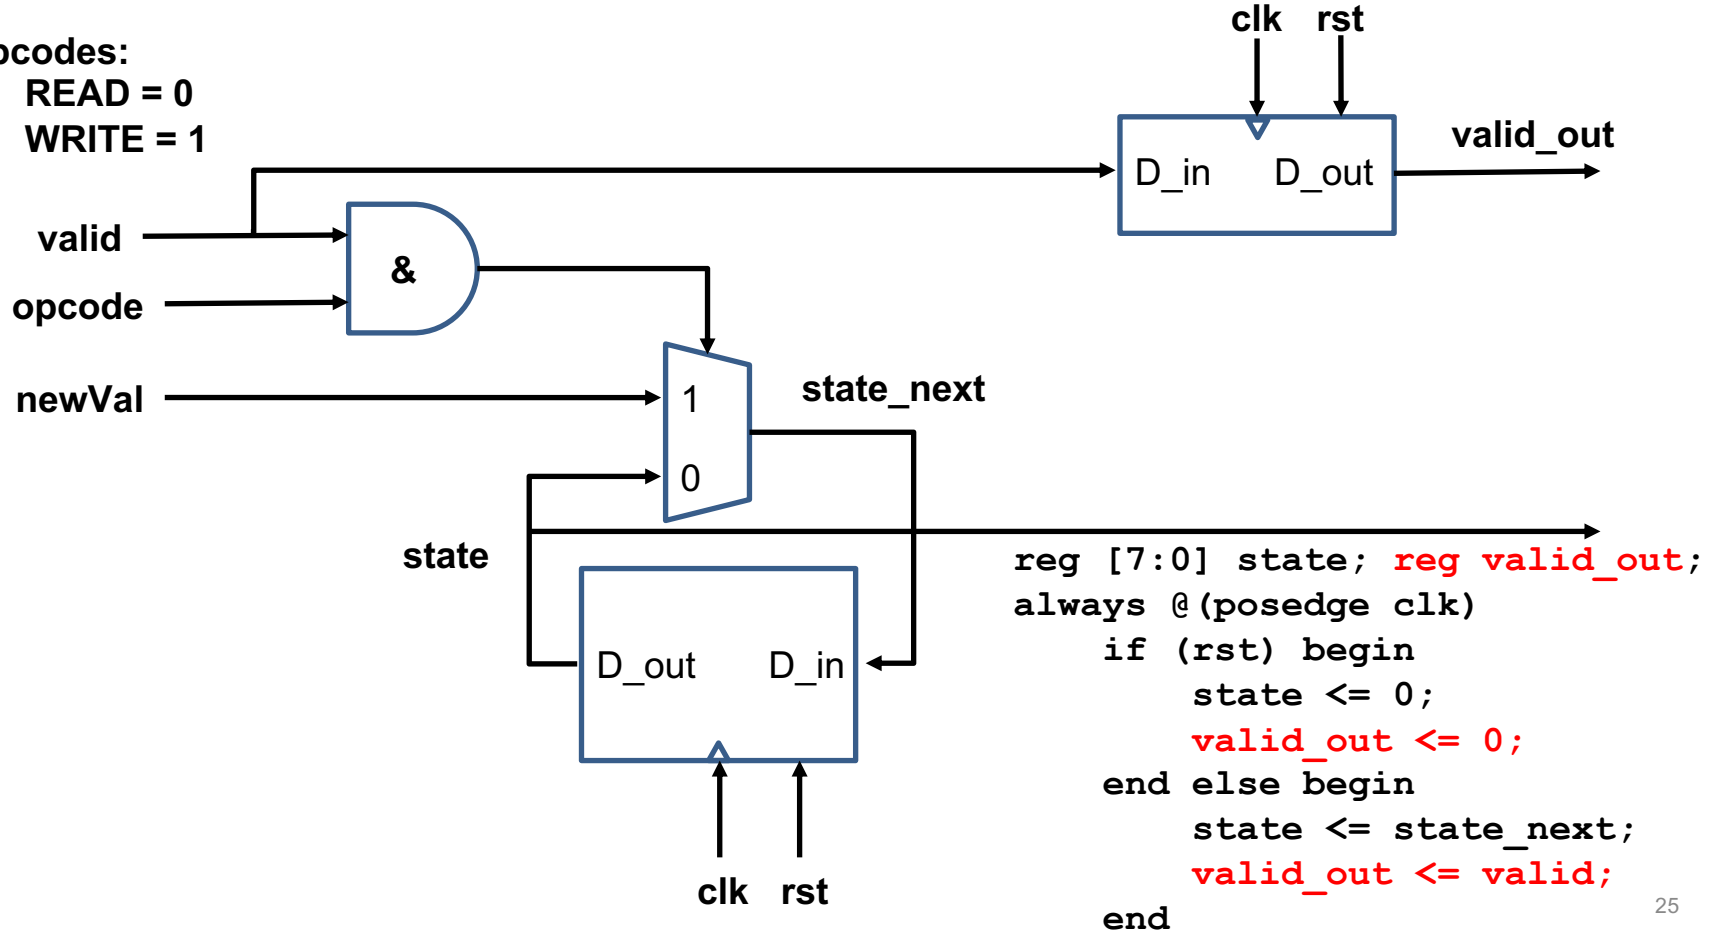
- READ = 0
- WRITE = 1



# Register the Outputs

opcodes:

- READ = 0
- WRITE = 1






# More Advanced State Machines

```
localparam STATE_1 = 0;
localparam STATE_2 = 1;
reg state, state_next;
reg [1:0] output_1;

always @(*) begin
    // defaults
    state_next = state;
    output_1 = 1;
    case (state)
        STATE_1: begin
            output_1 = 1;
            state_next = STATE_2;
        end

        STATE_2: begin
            output_1 = 2;
            state_next = STATE_1;
        end
    endcase
end

always @(posedge clk) begin
    if (rst)
        state <= STATE_1;
    else
        state <= state_next;
    end
end
```



# SDNet Extern API

## P4 API:

```
#define REG_READ 0
#define REG_WRITE 1
@Xilinx_MaxLatency(1)
@Xilinx_ControlWidth(0)
extern void <reg_name>_reg_srw<D>(in D newVal,
                                   in bit opCode,
                                   out D result);
```

## HDL Interface:

```
module my_reg_srw (
    input                clk,
    input                rst,
    input                input_VALID,
    input    [REG_WIDTH+OP_WIDTH:0] input_DATA,
    output               output_VALID,
    output    [REG_WIDTH-1:0] output_DATA
);

wire valid, stateful_valid, opcode;
wire [REG_WIDTH-1:0] newVal;
assign valid = input_VALID;
assign {stateful_valid, newVal, opcode} = input_DATA;
```

# Stateful\_Valid Signal

---

```
bit<8> result;  
if (p.hdr.invoke == 1) {  
    myReg_reg_srw(0, REG_WRITE, result);  
} else {  
    result = 32;  
}
```

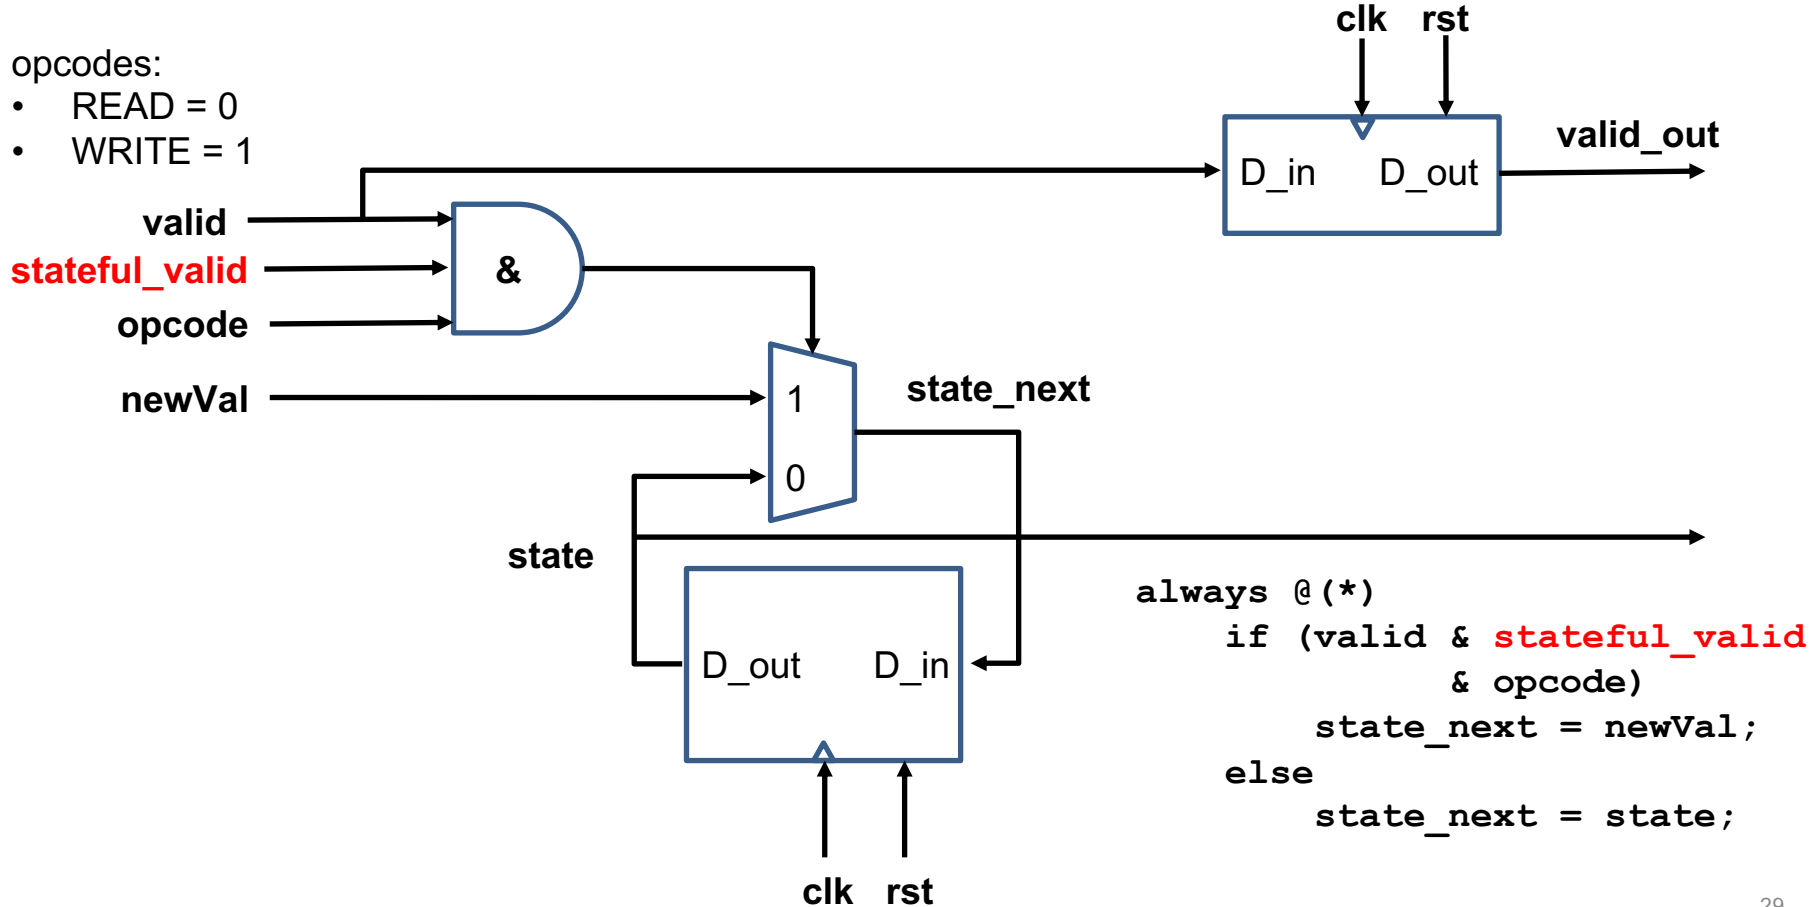
- **Two cases:**

- `p.hdr.invoke == 1` → `stateful_valid` signal will be set
- `p.hdr.invoke != 1` → `stateful_valid` signal will ***not*** be set
- **`input_VALID` signal *will* be asserted in both cases**

# State Machine

opcodes:

- READ = 0
- WRITE = 1



# SDNet Extern API

## HDL Interface:

```
module @MODULE_NAME@#(  
    parameter REG_WIDTH = @REG_WIDTH@,  
    parameter OP_WIDTH = 1)  
(  
    input                                clk_lookup,  
    input                                rst,  
    input                                tuple_in_@EXTERN_NAME@_input_VALID,  
    input [REG_WIDTH+OP_WIDTH:0] tuple_in_@EXTERN_NAME@_input_DATA,  
    output                                tuple_out_@EXTERN_NAME@_output_VALID,  
    output [REG_WIDTH-1:0] tuple_out_@EXTERN_NAME@_output_DATA  
);  
  
wire valid, stateful_valid, opcode;  
wire [REG_WIDTH-1:0] newVal;  
assign valid = tuple_in_@EXTERN_NAME@_input_VALID;  
assign {stateful_valid, newVal, opcode} = tuple_in_@EXTERN_NAME@_input_DATA;
```

# Adding extern support to P4→NetFPGA

- Update file: `$SUME_SDNET/bin/extern_data.py`

```
extern_data = {  
    ...  
    "reg_srw" : {"hdl_template_file": "externs/path/to/EXTERN_reg_srw_template.v",  
                 "replacements": {"@EXTERN_NAME@" : "extern_name",  
                                   "@MODULE_NAME@" : "module_name",  
                                   "@REG_WIDTH@" : "input_width(newVal)"}},  
    ...  
}
```

- **Commands used:**

- **extern\_name** – full name of extern function, determined by SDNet
- **module\_name** – name of the top level extern module, determined by SDNet
- **input\_width(field)** – width in bits of an input field, determined by P4 programmer

# Putting it all together: EXTERN\_reg\_srw\_template.v

```
module @MODULE_NAME@#(
    parameter REG_WIDTH = @REG_WIDTH@,
    parameter OP_WIDTH = 1)
(
    input                clk_lookup,
    input                rst,
    input                tuple_in_@EXTERN_NAME@_input_VALID,
    input [REG_WIDTH+OP_WIDTH:0] tuple_in_@EXTERN_NAME@_input_DATA,
    output               tuple_out_@EXTERN_NAME@_output_VALID,
    output [REG_WIDTH-1:0] tuple_out_@EXTERN_NAME@_output_DATA
);
```

```
// wire and reg declarations
wire valid, stateful_valid, opcode;
wire [REG_WIDTH-1:0] newVal;
reg [REG_WIDTH-1:0] state, state_next;
reg valid_out;
```

```
// decoding the inputs
assign valid = tuple_in_@EXTERN_NAME@_input_VALID;
assign {stateful_valid, newVal, opcode}
    = tuple_in_@EXTERN_NAME@_input_DATA;
```

```
// next state logic
always @(*)
    if (valid & stateful_valid & opcode)
        state_next = newVal;
    else
        state_next = state;
```

```
// state update / output logic
always @(posedge clk_lookup)
    if (rst) begin
        state <= 0;
        valid_out <= 0;
    end
    else begin
        state <= state_next;
        valid_out <= valid;
    end
end
```

```
// wire up the outputs
assign tuple_out_@EXTERN_NAME@_output_VALID = valid_out;
assign tuple_out_@EXTERN_NAME@_output_DATA = state;

endmodule
```

# Using our new extern: `srw_test.p4`

```
// extern declaration
#define REG_READ 0
#define REG_WRITE 1
@Xilinx_MaxLatency(1)
@Xilinx_ControlWidth(0)
extern void myReg_reg_srw(in bit<8> newVal,
                        in bit opCode,
                        out bit<8> result);

// match-action pipeline
control TopPipe(...) {

    apply {
        . . .
        bit<16> newVal;
        bit opcode;
        if (p.ethernet.etherType > 10) {
            newVal = p.ethernet.etherType;
            opcode = REG_WRITE;
        } else {
            newVal = 0; // unused
            opcode = REG_READ;
        }
        myReg_reg_srw(newVal, opcode, p.ethernet.etherType);
        . . .
    }
}
```

- [extern-tutorial branch:](#)
  - `srw_test` project



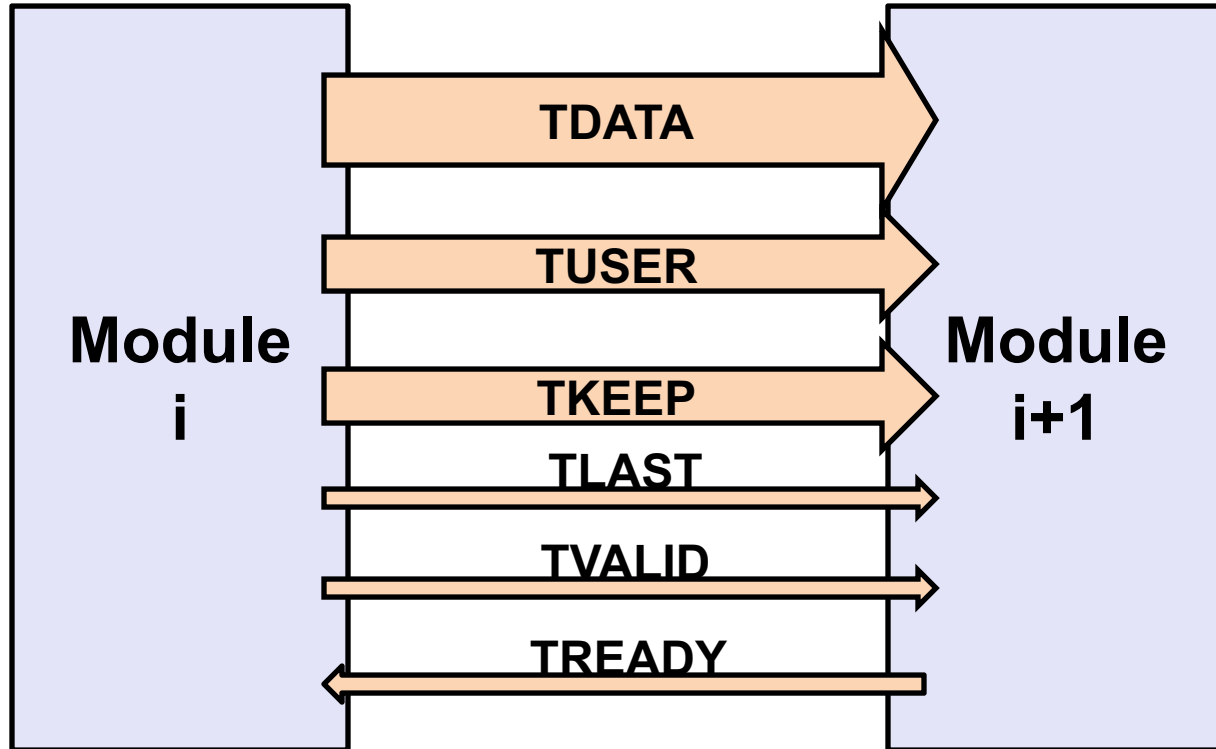
# What we didn't cover

---

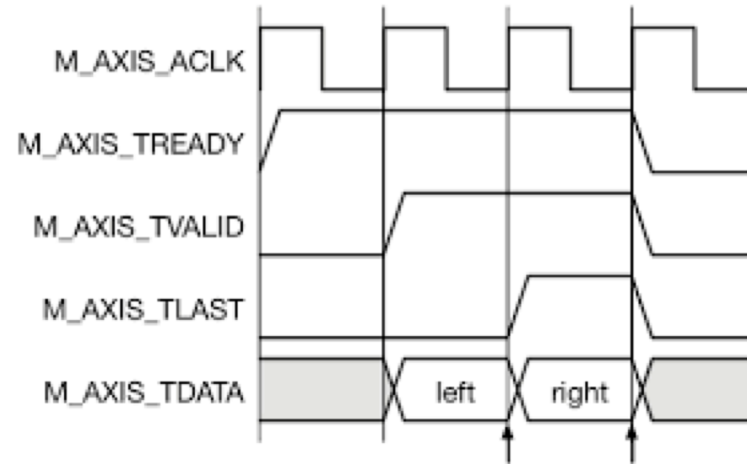
- Externs with control-plane interface
- BRAM based stateful extern
- Extern C++ implementations

# AXI4 Stream Interface

- Standardized interface for streaming packets between modules



# AXI4-Stream Handshake



- **TVALID & TREADY** → data is being transferred
- **TVALID & TREADY & TLAST** → the final word of the pkt is being transferred

# HDL Ethernet Parser

```
always @(*) begin
    // default values
    src_mac_w      = 0;
    dst_mac_w      = 0;
    eth_done_w     = 0;
    src_port_w     = 0;
    state_next     = state;

    case(state)
        /* read the input source header and get the first word */
        READ_MAC_ADDRESSES: begin
            if(valid) begin
                src_port_w  = tuser[SRC_PORT_POS+7:SRC_PORT_POS];
                dst_mac_w   = tdata[47:0];
                src_mac_w   = tdata[95:48];
                eth_done_w  = 1;
                state_next = WAIT_EOP;
            end
        end // case: READ_WORD_1

        WAIT_EOP: begin
            if(valid && tlast)
                state_next = READ_MAC_ADDRESSES;
            end
        endcase // case(state)
    end // always @ (*)
```

```
always @(posedge clk) begin
    if(reset) begin
        src_port <= {NUM_QUEUES{1'b0}};
        dst_mac  <= 48'b0;
        src_mac  <= 48'b0;
        eth_done <= 0;
        state    <= READ_MAC_ADDRESSES;
    end
    else begin
        src_port <= src_port_w;
        dst_mac  <= dst_mac_w;
        src_mac  <= src_mac_w;
        eth_done <= eth_done_w;
        state    <= state_next;
    end
end
```

# Parser Comparison

## Verilog

```
always @(*) begin
    src_mac_w    = 0;
    dst_mac_w    = 0;
    eth_done_w   = 0;
    src_port_w   = 0;
    state_next   = state;

    case(state)
        /* read the input source header and get the first word */
        READ_MAC_ADDRESSES: begin
            if(valid) begin
                src_port_w = tuser[Src_PORT_POS+7:Src_PORT_POS];
                dst_mac_w  = tdata[47:0];
                src_mac_w  = tdata[95:48];
                eth_done_w = 1;
                state_next = WAIT_EOP;
            end
        end // case: READ_WORD_1

        WAIT_EOP: begin
            if(valid && tlast)
                state_next = READ_MAC_ADDRESSES;
            end
        endcase // case(state)
    end // always @ (*)

always @(posedge clk) begin
    if(reset) begin
        src_port <= {NUM_QUEUES{1'b0}};
        dst_mac  <= 48'b0;
        src_mac  <= 48'b0;
        eth_done <= 0;
        state    <= READ_MAC_ADDRESSES;
    end
    else begin
        src_port <= src_port_w;
        dst_mac  <= dst_mac_w;
        src_mac  <= src_mac_w;
        eth_done <= eth_done_w;
        state    <= state_next;
    end // else: !if(reset)
end // always @ (posedge clk)
```

## P4

```
header Ethernet_h {
    bit<48> dstAddr;
    bit<48> srcAddr;
    bit<16> etherType;
}

parser TopParser(packet_in b,
                  out Parsed_packet p) {
    state start {
        b.extract(p.ethernet);
        transition accept;
    }
}
```

# Design Challenge Projects

- > **Mini research projects**
- > **Final presentation and demonstration (Wed June 12<sup>th</sup> 3:30-6:30pm)**
- > **Last year's projects:**
  - >> NetCache on NetFPGA
  - >> In-Network compression
  - >> Network accelerated sorting
  - >> P4 performance measurement system
  - >> INT & network cookies
  - >> Count Min Sketch for packet scheduling
- > **Proposal due Thursday (4/25)**
  - >> Project title
  - >> Problem statement and motivation
  - >> Initial thoughts for system design and final demo
- > **Project suggestions: Event-Driven P4 programs**

# Modern Programmable Data-Planes are Limited

## Two Key Limitations:

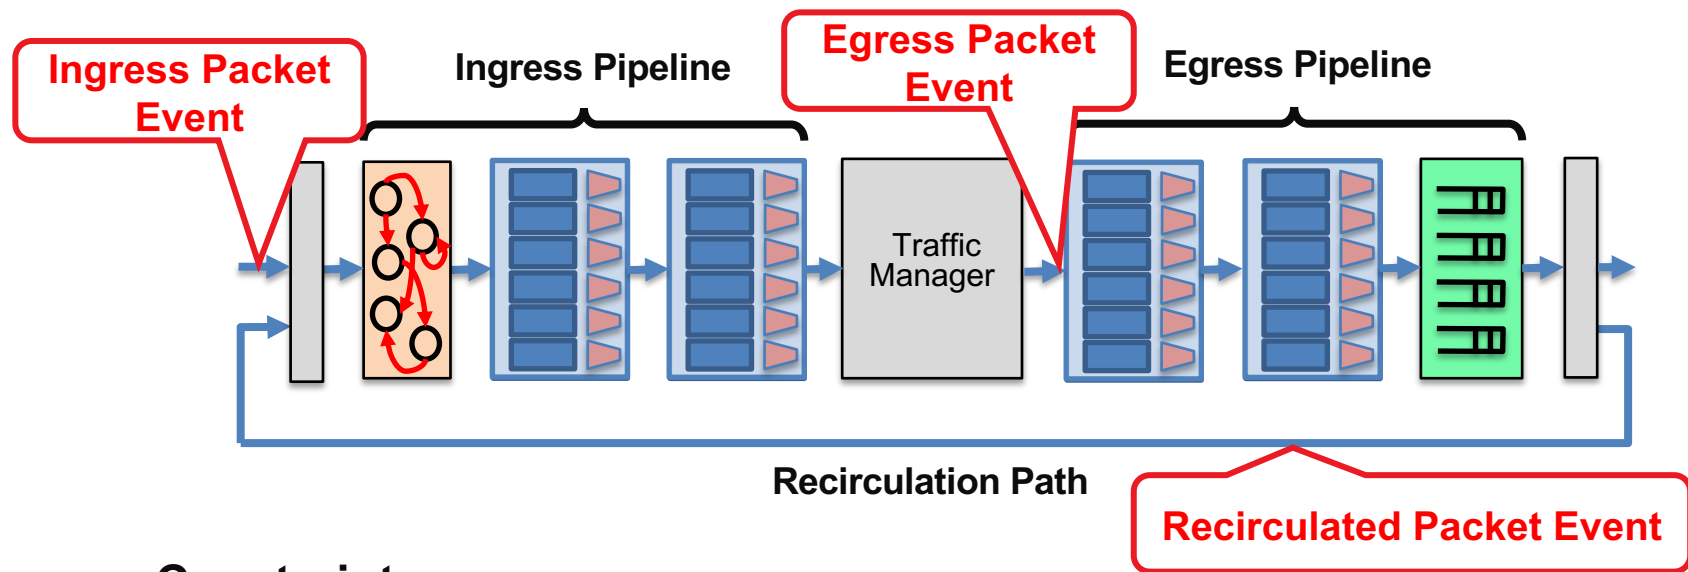
1. Deriving congestion signals for AQM, congestion control, load balancing, and telemetry applications

Common Congestion Signals	Other Congestion Signals
<ul style="list-style-type: none"><li>• Queue length</li><li>• Queue service rate</li><li>• Queueing delay</li></ul>	<ul style="list-style-type: none"><li>• Packet loss volume</li><li>• Rate of change of queue size</li><li>• Timestamp of buffer overflow/underflow events</li><li>• Per-active-flow buffer occupancy</li><li>• Etc...</li></ul>

2. Performing periodic tasks

- Periodic state update: Count-Min-Sketch, NumFabric, sPERC, AFD, etc.
- Periodic packet generation: HULA, BFD, etc.

# Modern PISA Switch



## Constraints:

- > No shared state across pipeline stages
- > Limited pipeline bandwidth
- > Limited pipeline stages
- > Limited memory
- > Limited per-packet computation
- > Limited control-plane bandwidth

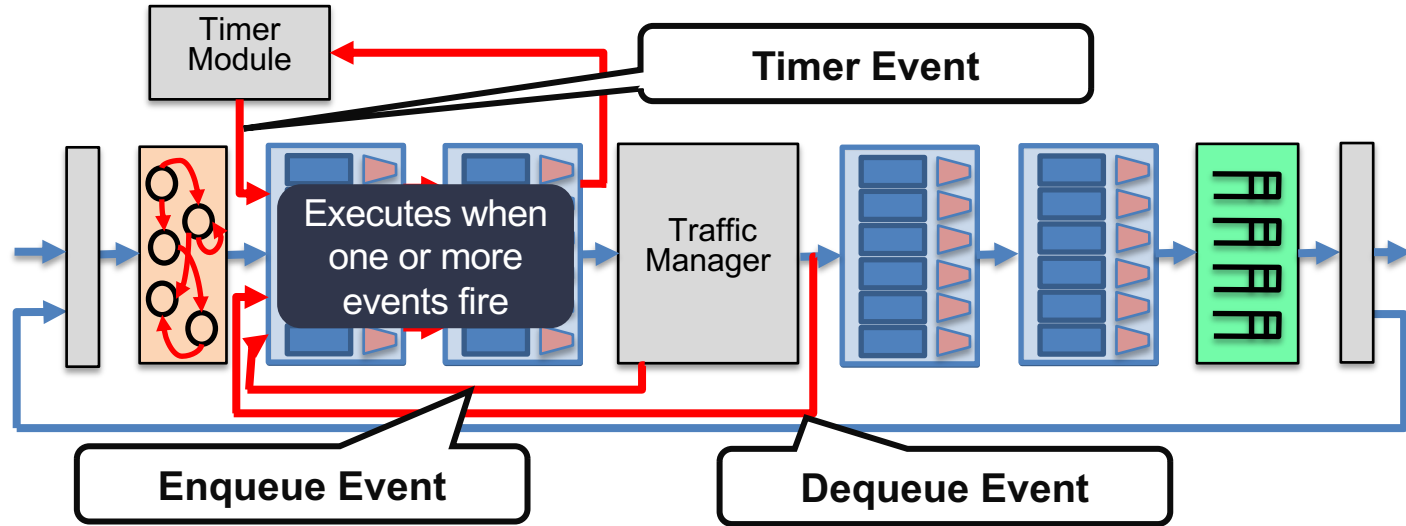
> **Only supports 3 event types**



# Data-Plane Events

		Event Type	Description
Packet & Metadata Events	{	Ingress Packet	Packet arrival
		Egress Packet	Packet departure
		Recirculated packet	Packet sent back to ingress
Metadata Events	{	Buffer Enqueue	Packet enqueued in buffer
		Buffer Dequeue	Packet dequeued from buffer
		Buffer Overflow / Packet Drop	Packet dropped at buffer
		Buffer Underflow	Buffer becomes empty
		Timer Event	Configurable timer expires
		Control-plane triggered	Control-plane triggers processing logic in data-plane
		Link Status Change	Link goes down / comes up

# Example Event-Driven Architecture

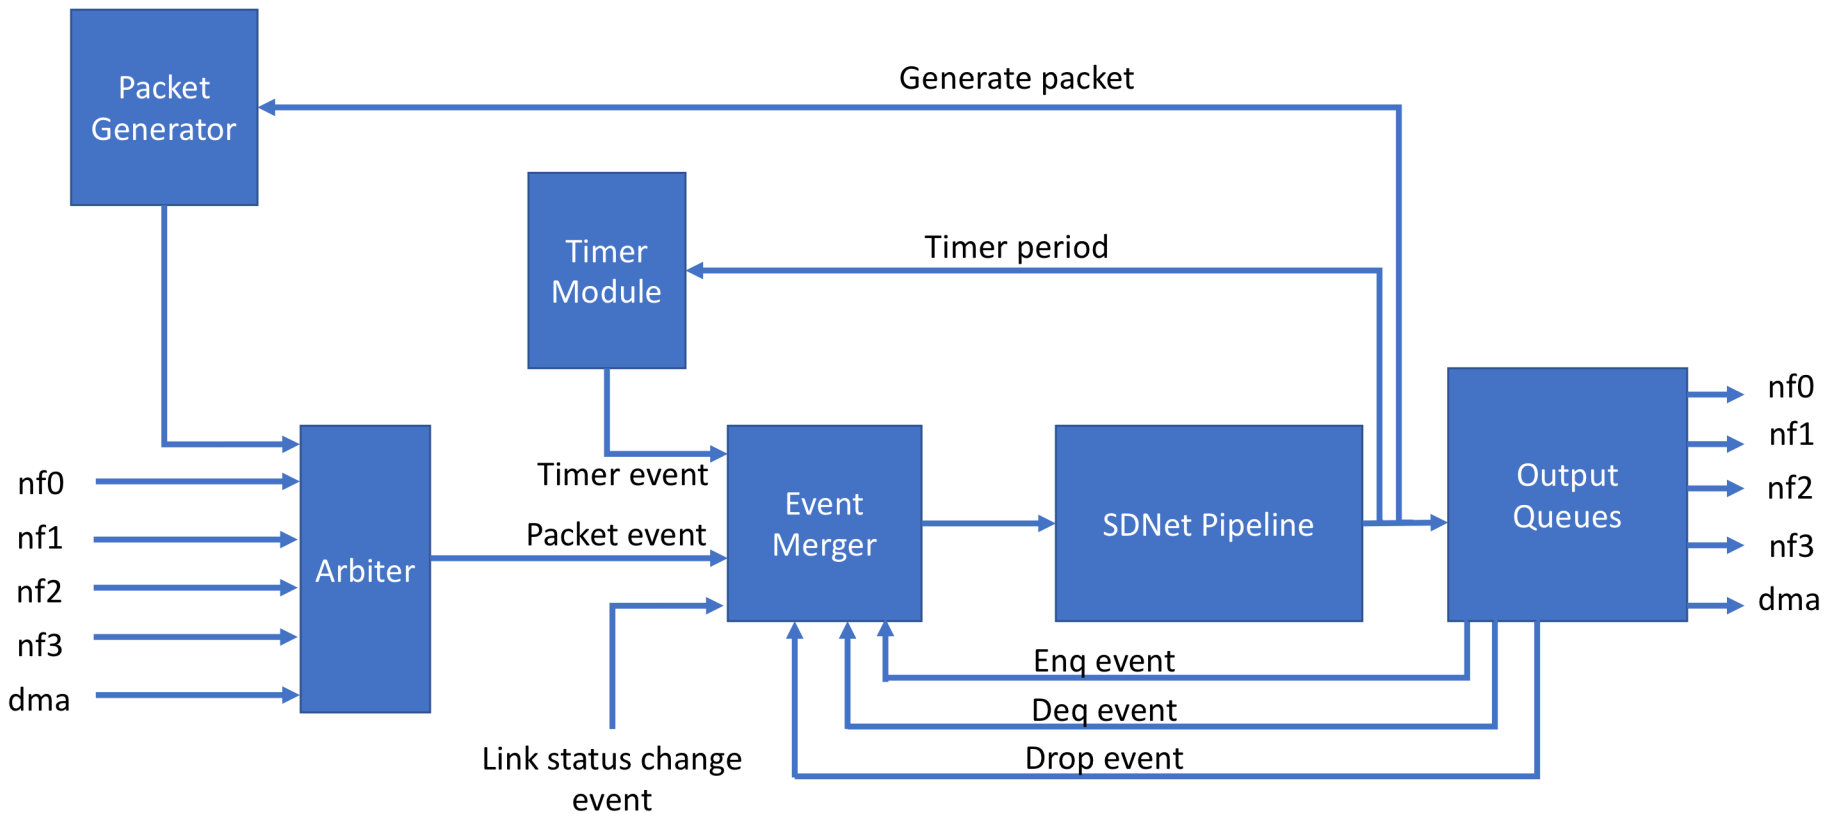


```
struct std_timer_t {  
    bool trigger;  
    bit<32> period;  
}
```

```
struct std_enq_t {  
    bool trigger;  
}  
struct enq_meta_t {  
    bit<8> flowID;  
    bit<32> pkt_len;  
}
```

```
struct std_deq_t {  
    bool trigger;  
}  
struct deq_meta_t {  
    bit<8> flowID;  
    bit<32> pkt_len;  
}
```

# SUME Event Switch Architecture



# SUME Event Switch Metadata

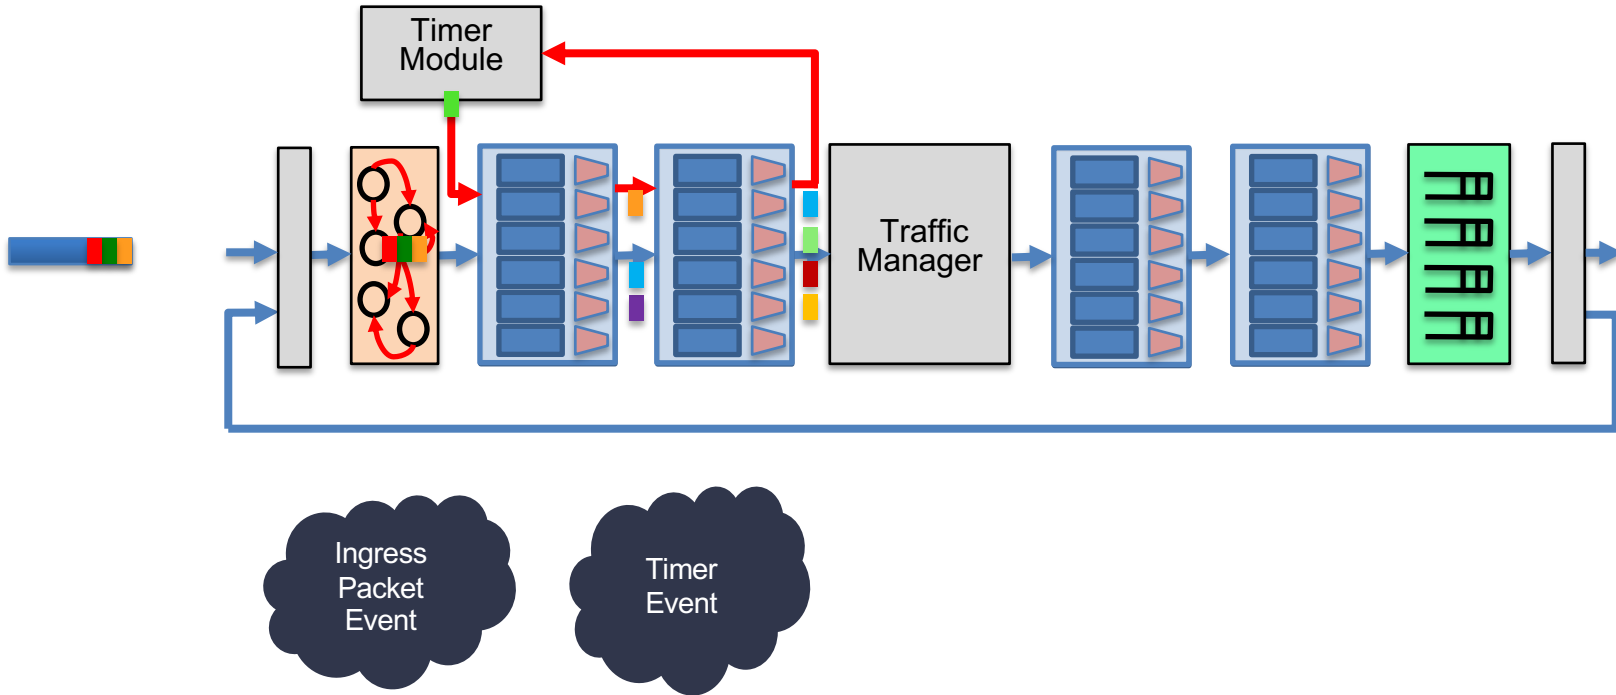
```
struct sume_metadata_t {  
    // request a packet to be generated  
    bit<1> gen_packet;  
    // which events have fired  
    bit<1> link_trigger;  
    bit<1> timer_trigger;  
    bit<1> drop_trigger;  
    bit<1> deq_trigger;  
    bit<1> enq_trigger;  
    bit<1> pkt_trigger;  
    // Current link status (one-hot)  
    bit<4> link_status;  
    // last timer event (20ns resolution)  
    bit<48> timer_now;  
    // timer event period (20ns resolution)  
    bit<32> timer_period;  
    // ports and metadata for events  
    bit<8> drop_port;  
    bit<8> deq_port;  
    bit<8> enq_port;  
    bit<32> drop_data;  
    bit<32> deq_data;  
    bit<32> enq_data;  
    // standard metadata fields  
    bit<8> dst_port;  
    bit<8> src_port;  
    bit<16> pkt_len;  
}
```

# Project Area Suggestions

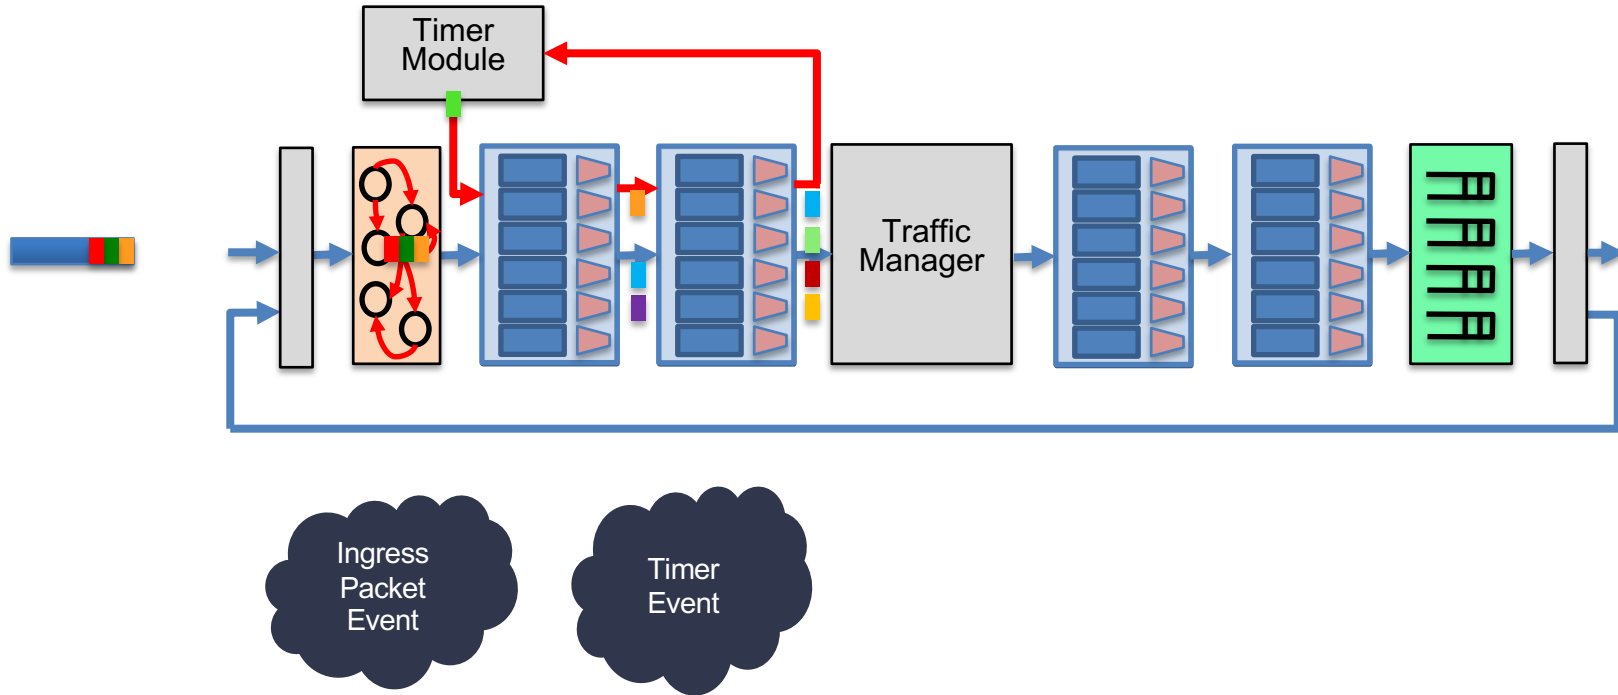
- > **Microburst Detection**
- > **Fast Reroute**
- > **Load Balancing**
- > **Network Monitoring**
- > **Active Queue Management**

**Extra Slides**

# Example Event-Driven Architecture



# Example Event-Driven Architecture





# Example Event-Driven Architecture

